

[\[SQL调优\] “查询SQL过滤和排序条件涉及的表字段未创建索引”引起慢查询问题，优化后执行时间从50+s下降到2s以下](#)

博客分类: [MySQL](#) [性能调优](#)

MySQL

前几天发现，线上portal“策略中心”的“证据管理”页面加载很慢。经排查发现，是由于 **riskbase_core** 库的 **evidence** 表未对 **gmt_create** 创建索引 和 **evidence_details** 表未对 **refuuid** 创建索引引起（因为查询条件涉及到这些字段），导致SQL执行时间要 1分钟+。

查询SQL

```
select d.type,d.value,e.fraud_type,e.evidence_time,e.evidence_origin,d.uuid,d.refuuid from evidence_details d inner join evidence e
on d.refuuid=e.uuid order by gmt_create desc limit 0,10;
```

经验教训

只要涉及到SQL查询条件（WHERE、ORDER BY）的相关字段，都应建立索引（唯一索引、联合索引）。

分析过程

1. 核实“SQL的执行时间”

```
select d.type,d.value,e.fraud_type,e.evidence_time,e.evidence_origin,d.uuid,d.refuuid from evidence_details d inner join evidence e
on d.refuuid=e.uuid order by gmt_create desc limit 0,10;
```

.....

10 rows in set (**1 min 11.27 sec**)

哇哦，该条SQL执行时间竟然需要1分11秒，太恐怖啦！！

2. 分析“该条SQL的查询执行计划”

```
explain select d.type,d.value,e.fraud_type,e.evidence_time,e.evidence_origin,d.uuid,d.refuuid from evidence_details d inner join
evidence e on d.refuuid=e.uuid order by gmt_create desc limit 0,10;
```

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	d	ALL	NULL	NULL	NULL	NULL	6452641	Using where; Using temporary; Using filesort
1	SIMPLE	e	eq_ref	uuid	uuid	96	riskbase_core.d.refuuid	1	NULL

从上面可以看出，查询第一步使用全表扫描（**ALL**），还涉及到临时表和文件排序（**Using where; Using temporary; Using filesort**）。所以，为了提高查询速度，尽量针对相关查询字段（``evidence_details`.`refuuid``、``evidence`.`uuid``、``evidence`.`gmt_create``）建立合理的索引。

3. 查看相应的索引是否创建

```
show index from `riskbase_core`.`evidence`;
```

Table	Non_unique	Key_name	Seq_in_index	Column_name	Collation	Cardinality	Sub_part	Packed	Null	Index_type	Comment	Index_comment
-------	------------	----------	--------------	-------------	-----------	-------------	----------	--------	------	------------	---------	---------------

```
| evidence |      0 | PRIMARY |      1 | id      | A      | 6471176 | NULL | NULL | | BTREE | | |
| evidence |      0 | uuid    |      1 | uuid    | A      | 6471176 | NULL | NULL | | BTREE | | |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
```

从上面看，`evidence`.`uuid` 字段的索引已创建。还需要创建`evidence`.`gmt_create`字段的索引，创建语句如下：

```
CREATE INDEX `dex_time` on `riskbase_core`.`evidence` (`gmt_create`);
```

```
show index from `riskbase_core`.`evidence_details`;
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| Table      | Non_unique | Key_name | Seq_in_index | Column_name | Collation | Cardinality | Sub_part | Packed | Null |
| Index_type | Comment    | Index_comment |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| evidence_details |      0 | PRIMARY |      1 | id      | A      | 6452885 | NULL | NULL | | BTREE | | |
|
| evidence_details |      0 | uuid    |      1 | uuid    | A      | 6452885 | NULL | NULL | YES | BTREE | | |
|
| evidence_details |      1 | idx_value |      1 | value   | A      | 6452885 | NULL | NULL | | BTREE | | |
|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
```

从上面看，未对`evidence_details`.`refuuid`字段创建索引。所以，需要对`evidence_details`.`refuuid`字段创建索引，创建语句如下：

```
CREATE INDEX `dex_uuid` on `riskbase_core`.`evidence_details` (`refuuid`);
```

4. 再次查看“该SQL的查询执行计划”

```
explain select d.type,d.value,e.fraud_type,e.evidence_time,e.evidence_origin,d.uuid,d.refuuid from evidence_details d inner join
evidence e on d.refuuid=e.uuid order by gmt_create desc limit 0,10;
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table | type | possible_keys | key      | key_len | ref      | rows | Extra |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 1 | SIMPLE      | e      | index | index | uuid    | 6        | NULL     | 10 | NULL |
| 1 | SIMPLE      | d      | ref   | dex_uuid | dex_uuid | 97       | riskbase_core.e.uuid | 1 | NULL |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
```

从SQL查询执行计划看，相关索引都使用上了，看起来应该没问题了。现在就用SQL语句测试一下执行时间吧

5. 再次执行该SQL，用“执行时间”来验证优化效果

```
select d.type,d.value,e.fraud_type,e.evidence_time,e.evidence_origin,d.uuid,d.refuuid from evidence_details d inner join evidence e
on d.refuuid=e.uuid order by gmt_create desc limit 0,10; .....
```

10 rows in set (0.00 sec)

好了，从“执行时间”看，问题已彻底修复了。

6. 看一下优化前后，页面的加载效果

优化前 **57.41s**

方法	文件	域名	类型	已传输	大小	0 毫秒	10.24 秒	20.48 秒	30.72 秒	40.96 秒	51.20 秒	
304 GET	switch2.js?v=1432026875978	portal.fraudmetrix.cn	js	4.76 KB	4.40 KB							→ 329 ms
304 GET	search.js?v=1432026875978	portal.fraudmetrix.cn	js	1.91 KB	1.91 KB							→ 338 ms
304 GET	jquery.placeholder.min.js?v=1432026875978	portal.fraudmetrix.cn	js	2.11 KB	2.11 KB							→ 389 ms
304 GET	ui-select.js?v=1432026875978	portal.fraudmetrix.cn	js	15.64 KB	15.05 KB							→ 360 ms
304 GET	jquery-ui.js?v=1432026875978	portal.fraudmetrix.cn	js	425.63 KB	425.63 KB							→ 357 ms
304 GET	pages.js?v=1432026875978	portal.fraudmetrix.cn	js	20.85 KB	19.60 KB							→ 361 ms
304 GET	jquery.popup.js?v=1432026875978	portal.fraudmetrix.cn	js	2.30 KB	2.28 KB							→ 379 ms
304 GET	WdatePicker.js?v=1432026875978	portal.fraudmetrix.cn	js	9.72 KB	9.72 KB							→ 302 ms
304 GET	pop-mag.js?v=1432026875978	portal.fraudmetrix.cn	js	20.43 KB	19.89 KB							→ 414 ms
304 GET	toolTip.js?v=1432026875978	portal.fraudmetrix.cn	js	10.55 KB	10.37 KB							→ 413 ms
304 GET	jquery.tdValidator.min.js?v=1432026875978	portal.fraudmetrix.cn	js	18.95 KB	16.48 KB							→ 426 ms
304 GET	jquery.nicescroll.min.js?v=1432026875978	portal.fraudmetrix.cn	js	55.54 KB	55.54 KB							→ 426 ms
200 GET	hm.js?67c1bdc2a820e48e03e1a70e0f371	hm.baidu.com	js	7.59 KB	19.56 KB							→ 307 ms
200 POST	PartnerAppSelect.json	portal.fraudmetrix.cn	html	1.77 KB	5.40 KB							→ 160 ms
200 POST	evidenceManage.json	portal.fraudmetrix.cn	html	0.13 KB	0.11 KB							→ 1625 ms
200 POST	evidenceManage.json	portal.fraudmetrix.cn	html	1.23 KB	3.74 KB							→ 56 ms
200 GET	hm.gif?cc=1&cd=1&cl=24-bit&ds=1440x900&ep=...	hm.baidu.com	gif	0.04 KB	0.05 KB							→ 37 ms
200 GET	hm.gif?cc=1&cd=1&cl=24-bit&ds=1440x900&et=0...	log.hm.baidu.com	gif	0.04 KB	0.05 KB							→ 185 ms
200 GET	hm.gif?cc=1&cd=1&cl=24-bit&ds=1440x900&ep=...	log.hm.baidu.com	gif	0.04 KB	0.05 KB							→ 43 ms

优化后 1.79s

方法	文件	域名	类型	已传输	大小	0 毫秒	325 毫秒	640 毫秒	955 毫秒	1,270 毫秒	1,585 毫秒	1,900 毫秒	2,215 毫秒	2,530 毫秒	
304 GET	switch2.js?v=1432026875978	portal.fraudmetrix.cn	js	4.76 KB	4.40 KB										→ 291 ms
304 GET	search.js?v=1432026875978	portal.fraudmetrix.cn	js	1.91 KB	1.91 KB										→ 326 ms
304 GET	jquery.placeholder.min.js?v=1432026875978	portal.fraudmetrix.cn	js	2.11 KB	2.11 KB										→ 326 ms
304 GET	ui-select.js?v=1432026875978	portal.fraudmetrix.cn	js	15.64 KB	15.05 KB										→ 329 ms
304 GET	jquery-ui.js?v=1432026875978	portal.fraudmetrix.cn	js	425.63 KB	425.63 KB										→ 320 ms
304 GET	pages.js?v=1432026875978	portal.fraudmetrix.cn	js	20.85 KB	19.60 KB										→ 330 ms
304 GET	jquery.popup.js?v=1432026875978	portal.fraudmetrix.cn	js	2.30 KB	2.28 KB										→ 320 ms
304 GET	WdatePicker.js?v=1432026875978	portal.fraudmetrix.cn	js	9.72 KB	9.72 KB										→ 301 ms
304 GET	pop-mag.js?v=1432026875978	portal.fraudmetrix.cn	js	20.43 KB	19.89 KB										→ 391 ms
304 GET	toolTip.js?v=1432026875978	portal.fraudmetrix.cn	js	10.55 KB	10.37 KB										→ 360 ms
304 GET	jquery.tdValidator.min.js?v=1432026875978	portal.fraudmetrix.cn	js	18.95 KB	16.48 KB										→ 360 ms
304 GET	jquery.nicescroll.min.js?v=1432026875978	portal.fraudmetrix.cn	js	55.54 KB	55.54 KB										→ 409 ms
200 GET	hm.js?67c1bdc2a820e48e03e1a70e0f371	hm.baidu.com	js	7.59 KB	19.56 KB										→ 305 ms
200 POST	PartnerAppSelect.json	portal.fraudmetrix.cn	html	1.77 KB	5.40 KB										→ 56 ms
200 POST	evidenceManage.json	portal.fraudmetrix.cn	html	0.13 KB	0.11 KB										→ 1677 ms
200 POST	evidenceManage.json	portal.fraudmetrix.cn	html	1.27 KB	3.74 KB										→ 55 ms
200 GET	hm.gif?cc=1&cd=1&cl=24-bit&ds=1440x900&ep=...	log.hm.baidu.com	gif	0.04 KB	0.05 KB										→ 142 ms
200 GET	hm.gif?cc=1&cd=1&cl=24-bit&ds=1440x900&et=0...	hm.baidu.com	gif	0.04 KB	0.05 KB										→ 118 ms
200 GET	hm.gif?cc=1&cd=1&cl=24-bit&ds=1440x900&ep=...	hm.baidu.com	gif	0.04 KB	0.05 KB										→ 39 ms

页面加载时间，从优化前的 57.41s 降低到 1.79s。

但可能细心的朋友已经发现，优化后 evidenceManage.json 的响应时间还需要 1.677s，而上面那条SQL的执行时间已经降低到10ms之内。那其它的时间都消耗在哪里啦？

其它故事

其实 evidenceManage.json 接口除了执行上面那条SQL语句，还执行了一条count(*)的SQL语句。而正是count(*) SQL语句耗时了 1.67s，哦哦，原来这里还有其它故事啊～～～

总行数统计

select count(*) from evidence_details

但通过对“count(*) SQL语句”的查询执行计划的分析，发现对于InnoDB引擎（[14.2 InnoDB Concepts and Architecture](#)），很难对此再优化了。

（建议：有count(*)相关的操作，数据表的存储引擎（ENGINE）尽量设计为MyISAM（[15.2 The MyISAM Storage Engine](#)），除非该表涉及事务操作！）

总结

针对 InnoDB 存储引擎：

- 索引（index）查询类型的查询要快于范围（range）查询类型
- 二级索引（dex_uuid）类型的查询要快于主键索引（PRIMARY）类型

参考

[\[MySQL FAQ系列\] 为何 InnoDB 表 select count\(*\) 很慢 -- 叶金荣\(yejr\)](#)

[\[InnoDB系列\] InnoDB 表如何更快得到 count\(*\) 结果 -- 叶金荣\(yejr\)](#)

分析过程

```
mysql> explain select count(*) from evidence_details \G
```

```
***** 1. row *****
```

```
id: 1
```

```
select_type: SIMPLE
```

```
table: evidence_details
```

```
type: index
```

```
possible_keys: NULL
```

```
key: dex_uuid
```

```
key_len: 97
```

```
ref: NULL
```

```
rows: 6479241
```

```
Extra: Using index
```

```
1 row in set (0.00 sec)
```

```
mysql> select count(*) from evidence_details \G
```

```
***** 1. row *****
```

```
count(*): 7640484
```

```
1 row in set (1.67 sec)
```

```
mysql> explain select count(*) from evidence_details where id >= 0 \G
```

```
***** 1. row *****
```

```
id: 1
```

```
select_type: SIMPLE
```

```
table: evidence_details
```

```
type: range
```

```
possible_keys: PRIMARY
```

```
key: PRIMARY
```

```
key_len: 8
```

```
ref: NULL
```

```
rows: 3239629
```

```
Extra: Using where; Using index
```

```
1 row in set (0.00 sec)
```

```
mysql> select count(*) from evidence_details where id >= 0 \G
```

```
***** 1. row *****
```

count(*): 7640505

1 row in set (2.51 sec)

mysql> explain select count(id) from evidence_details \G

***** 1. row *****

id: 1

select_type: SIMPLE

table: evidence_details

type: **index**

possible_keys: NULL

key: dex_uuid

key_len: 97

ref: NULL

rows: 6479287

Extra: **Using index**

1 row in set (0.00 sec)

mysql> select count(id) from evidence_details \G

***** 1. row *****

count(id): 7640530

1 row in set (1.83 sec)

mysql> explain select count(id) from evidence_details where id >= 0 \G

***** 1. row *****

id: 1

select_type: SIMPLE

table: evidence_details

type: range

possible_keys: **PRIMARY**

key: PRIMARY

key_len: 8

ref: NULL

rows: 3239652

Extra: Using where; Using index

1 row in set (0.00 sec)

mysql> select count(id) from evidence_details where id >= 0 \G

***** 1. row *****

count(id): 7640547

1 row in set (2.64 sec)

```
mysql> explain select count(`uuid`) from evidence_details \G
```

```
***** 1. row *****

id: 1

select_type: SIMPLE
table: evidence_details
type: index
possible_keys: NULL
key: uuid
key_len: 99
ref: NULL
rows: 6479323
Extra: Using index
1 row in set (0.00 sec)
```

```
mysql> select count(`uuid`) from evidence_details \G
```

```
***** 1. row *****

count(`uuid`): 7640564
1 row in set (2.20 sec)
```

```
mysql> explain select count(`uuid`) from evidence_details where id >= 0 \G
```

```
***** 1. row *****

id: 1

select_type: SIMPLE
table: evidence_details
type: range
possible_keys: PRIMARY
key: PRIMARY
key_len: 8
ref: NULL
rows: 3239705
Extra: Using where
1 row in set (0.00 sec)
```

```
mysql> select count(`uuid`) from evidence_details where id >= 0 \G
```

```
***** 1. row *****

count(`uuid`): 7640657
1 row in set (3.13 sec)
```

```
mysql> explain select count(`refuuid`) from evidence_details \G
```

```
***** 1. row *****
```

```
id: 1
```

```
select_type: SIMPLE
```

```
table: evidence_details
```

```
type: index
```

```
possible_keys: NULL
```

```
key: dex_uuid
```

```
key_len: 97
```

```
ref: NULL
```

```
rows: 6479554
```

```
Extra: Using index
```

```
1 row in set (0.00 sec)
```

```
mysql> select count(`refuuid`) from evidence_details \G
```

```
***** 1. row *****
```

```
count(`refuuid`): 7640812
```

```
1 row in set (2.09 sec)
```

```
mysql> explain select count(`refuuid`) from evidence_details where id >= 0 \G
```

```
***** 1. row *****
```

```
id: 1
```

```
select_type: SIMPLE
```

```
table: evidence_details
```

```
type: range
```

```
possible_keys: PRIMARY
```

```
key: PRIMARY
```

```
key_len: 8
```

```
ref: NULL
```

```
rows: 3239794
```

```
Extra: Using where
```

```
1 row in set (0.00 sec)
```

```
mysql> select count(`refuuid`) from evidence_details where id >= 0 \G
```

```
***** 1. row *****
```

```
count(`refuuid`): 7640852
```

```
1 row in set (3.31 sec)
```